

SOC Homelab — Master Documentation

Single, synthesized reference for the SOC Homelab project. Current state only — full history and evidence live in the per-topic source docs linked throughout and indexed at the bottom of this file.

Status labels:  verified/current ·  open item, needs action ·  planned, not built.

Table of contents

1. [Project overview & purpose](#)
 2. [Architecture & network design](#)
 3. [Systems in detail](#)
 4. [Architecture, security & AI rules](#)
 5. [Operational guides](#)
 6. [Detection use cases & incident response](#)
 7. [Troubleshooting history](#)
 8. [Project timeline](#)
 9. [Asset inventory](#)
 10. [Glossary](#)
 11. [Current project status / what's next](#)
 12. [Attack scope: agreed Red Team test plan](#)
-

1. Project overview & purpose

SOC Homelab (originally "Fortress Bazzite") is Joost Hebly's private cybersecurity training environment — a real SOC, built and operated end-to-end, including the real problems that come with it.


Focus areas: **Blue Team** (Security Onion/Kibana monitoring, incident response, detection engineering), **Red Team** (vulnerability testing, exploitation, AD attacks), realistic **Active Directory**, general **network defense**. Also a portfolio piece — Joost produces formal reports/documentation from it alongside the working lab.

Origin: started as Fortress Bazzite (self-built Suricata+Zeek+Wazuh stack, design doc 2026-07-05). In practice, **Security Onion** was chosen instead — same goal, one bundled platform (it already includes Suricata, Zeek, and Fleet/Elastic Agent in place of Wazuh). Original detection goals carried forward — [§6](#).

AI collaboration: ChatGPT for architecture/planning/strategy, Claude Code for terminal work/documentation/scripts/implementation. Both operate under the AI rules in [§4](#). **Golden rule:** no destructive changes without approval; all changes documented.

2. Architecture & network design

Network map



Security Onion has two NICs: `pentest-lab` (.30, Fleet/Kibana/SSH/web) and `monitor-net` (no IP, mirrored traffic only).

IP address table

IP	System	virsh VM name	SSH alias	SSH user	Status
192.168.50.1	OPNsense (firewall/gateway)	OPNsense-FW	opnsense	root	✓
192.168.50.10	DC01 (AD, PDC Emulator, domain pentest.lab)	DC01	dc01	Administrator	✓
192.168.50.20	WIN11-01 (domain-joined as DESKTOP- EFKB8GQ)	WIN11-01	win11-01	pentest\administrator (key auth confirmed working)	✓
192.168.50.30	Security Onion 3.1.0 (SIEM/IDS/Fleet)	SOC- SecurityOnion	security- onion	socadmin	✓
192.168.50.40	ubuntu-server-01 (live OWASP Juice Shop on :3000)	ubuntu-server- 01	ubuntu- server	sysadmin (key auth confirmed working)	✓
192.168.50.50	Kali Linux (Red Team)	ATTACK-Kali ⚠ <i>leading space in the name</i>	kali	blue1	✓
192.168.50.70	Metasploitable2 (vulnerable target)	Target- Metasploitable2	(none)	—	✓

Known naming bug: the VM name `ATTACK-Kali` has a literal leading space in libvirt (broke name-matching scripts once — [§7](#)). New scripts should match on UUID instead.

IP authority: `docs/ASSET_INVENTORY.md` and `NETWORK.md` (2026-07-13) are authoritative — older docs and the 2026-07-09 daily report contain now-superseded IPs for Security Onion, Kali, ubuntu-server-01, and Metasploitable2; not repeated here.

Virtual networks (libvirt)

Network	Purpose
pentest-lab	Main network, 192.168.50.0/24, isolated. No libvirt-native DHCP — OPNsense handles this exclusively.
monitor-net	Isolated, no IP addressing, exclusively for traffic mirroring to Security Onion.
default	Standard libvirt network, not used by lab VMs.

Physical host

Bazzite Linux. CPU: Intel Core i9-11900K (8c/1s) · GPU: NVIDIA RTX 3090 · RAM: 62 GiB · WiFi: Intel AX210 (Wi-Fi 6E) · Virtualization: KVM/QEMU/libvirt/virt-manager. All  live-verified.  WiFi PCI-passthrough to Kali and host-level Docker not reconfirmed — Juice Shop's actual location is `ubuntu-server-01`, not the host.

Traffic mirroring (`soc-mirror.service`)

Event-driven (`scripts/soc-mirror.sh` , triggered by a libvirt qemu hook on every VM start/stop) — no polling timer. When a `pentest-lab` VM starts, its traffic mirrors to Security Onion's `monitor-net` interface; stale rules clean up automatically on Security Onion restart. Status: `scripts/soc-mirror.sh --status` .

DNS and DHCP

DHCP: Kea DHCPv4 on OPNsense, LAN interface only. Subnet `192.168.50.0/24` , dynamic pool `.100 – .200` , 7 static reservations (one per lab VM by MAC — the canonical IP plan, see [§9](#)). Option 6 pushes `192.168.50.10` (DC01) as DNS server. Kea DHCPv6 exists but has no interface assigned (inactive). **DNS:** DC01 (AD DNS) authoritative for `pentest.lab` ; OPNsense's Unbound forwards the `pentest.lab` domain specifically to `192.168.50.10` and has one host override (`dc01.pentest.lab` → `.10`), everything else resolved/forwarded normally. Confirmed 2026-07-13 via a full read-only OPNsense audit — full detail: `docs/OPNSENSE_AUDIT_2026-07-13.md` .

Security principles

Segmentation (`monitor-net` fully separate from `pentest-lab`) · least privilege (Security Onion hostgroups) · monitor first, block second · every change documented.

Planned network improvements (not yet built)

VLAN segmentation · separate management/attack networks · additional Windows clients · honeypots · more SOC sensors.

3. Systems in detail

All facts below are live-verified 2026-07-13 unless marked otherwise.

3.1 Bazzite host & virtualization

Physical KVM/QEMU/libvirt/virt-manager host, doubling as Joost's daily desktop. Storage: qcow2 (snapshots, thin provisioning). Basic lifecycle:

```
virsh list --all
virsh start VMNAME
virsh shutdown VMNAME          # graceful; --destroy only if stuck
```

⚠ QEMU guest agent isn't configured by default per VM — install/enable in-guest before relying on guest-management features.

Own Elastic Agent (2026-07-14): the host itself (IP 192.168.50.254 on the virbr10 / pentest-lab bridge — distinct from OPNsense's .1) runs a log/metrics-only Elastic Agent (journald system.auth / system.syslog + system/metrics, no Elastic Defend by design — this is the machine every VM depends on). Confirmed Healthy/Connected in Fleet, confirmed to survive a full host reboot with zero manual steps. Ingest-side verified end-to-end the same day: a deliberate packet capture (zero TCP resets) plus a targeted Elasticsearch/Hunt query for the exact test window confirmed all 3 test sudo events fully indexed with correct field-level detail. An earlier same-day finding that journald logs never arrived was itself wrong — a diagnostic-method problem (the Fleet data-streams API doesn't reliably surface this bursty dataset), not a real delivery failure. Full investigation and methodology notes: docs/troubleshooting/08_bazzite_host_elastic_agent.md . Quick health check across this host + all 7 lab VMs: scripts/soc-health-check.sh .

3.2 OPNsense (firewall/gateway)

Central firewall/gateway for 192.168.50.0/24 — firewall, routing, DHCP, DNS forwarding. IP .1 . SSH alias opnsense , user root , password-only login by design (no key auth configured — confirmed via the OPNsense audit, not a regression). OPNsense 26.1.11_6-amd64 / FreeBSD 14.3. No custom firewall rules beyond the LAN/WAN defaults (no inter-VM segmentation at this layer), no VPN configured, single local admin user. Full configuration audit: docs/OPNSENSE_AUDIT_2026-07-13.md .

3.3 DC01 — Active Directory Domain Controller

Windows Server 2022 Standard Evaluation. Domain pentest.lab , DC01 is PDC Emulator (only DC). IP .10 . SSH alias dc01 / Administrator . Timezone W. Europe Standard Time (cosmetic; UTC clock is authoritative).

AD structure is built (not just planned):

```
DC=pentest,DC=lab
├─ OU=Domain Controllers      (DC01)
├─ OU=Admins                  (IT Admin 01)
├─ OU=AD-Users                (soctest, Helpdesk 01, Employee 01, Manager 01, HR
01, Finance 01)
├─ OU=Workstations            (empty)
├─ OU=Servers                 (empty)
├─ OU=Groups                  (SOC-Analysts: soctest · Helpdesk: no members)
└─ OU=Service-Accounts        (SQL Service)
```

⚠ **Known gaps** (real, deliberately left as-is pending §12): `OU=Workstations / OU=Servers` empty (WIN11-01 not moved in); `IT Admin 01` has no elevated group membership despite its name/OU; `Helpdesk` group has no members; role accounts (`Employee/Manager/HR/Finance 01`) are undifferentiated `Domain Users`.

Monitoring: Elastic Agent + Sysmon 15.21 (SwiftOnSecurity config) enrolled in Fleet, **Healthy** — full root-cause history in §7.6. NTP via `pool.ntp.org`; `vmictimesync` disabled (was fighting NTP at boot).

3.4 WIN11-01 — Windows 11 workstation

IP `.20`, SSH alias `win11-01` (`pentest\administrator`, key auth confirmed working — see `docs/troubleshooting/09_win11-01_ssh_access.md`). Domain-joined as `DESKTOP-EFKB8GQ` (never renamed, still in default `Computers` container). As of 2026-07-13, only TCP 135 was reachable; SMB/RDP/WinRM/139 were blocked by Windows Firewall. As of 2026-07-14, port 22 (SSH) is also open — Joost enabled OpenSSH Server himself via the VM console, a deliberate change, not a regression. Other ports not re-verified this session, assumed unchanged. ⚠ Intended training purpose still undecided beyond "becomes a target after cleanup" — §12.

Monitoring (2026-07-14): Elastic Agent 9.3.3 + Sysmon 15.21 (SwiftOnSecurity config) enrolled in Fleet under the same `endpoints-initial` policy DC01 uses (Windows Event Logs, Sysmon, Elastic Defend, osquery, metrics). **Healthy** — reached that state after Fleet's server-side view lagged the agent's own local status for ~12 minutes post-enrollment, comparable to DC01's documented stabilization delay, not a stuck state. Confirmed in Hunt: ~6,351 total events, ~1,004 `windows.sysmon_operational` events, current activity. Full investigation, including a firewall hostgroup step that was tested and found to be a no-op (not the fix) and a `wsasend` connection-reset message shown to also occur on DC01's healthy agent (so not diagnostic of a fault): `docs/troubleshooting/10_win11-01_sysmon_elastic_agent.md`.

3.5 ubuntu-server-01 — Linux server, active Red Team target

IP `.40` — confirmed via OPNsense's own Kea DHCP reservation database. Previously seen drifting to `.100` (dynamic pool) after a reboot; **root cause proven and fixed 2026-07-14** (`docs/troubleshooting/12_ubuntu-server-01_dhcp_reservation_fix.md`) — this Ubuntu image runs two DHCP negotiations per boot (an early dracut-fallback one, then the real netplan one), and without `dhcp-identifier: mac` the second used a non-MAC client identifier that Kea's reservation didn't match. Fixed with one netplan line; validated across a full reboot. SSH alias `ubuntu-server / sysadmin` (not `ubuntu` as earlier documented); key auth confirmed working. Elastic Agent installed and Healthy in Fleet (log/metrics-only, same scope as the Bazzite host) — `docs/troubleshooting/11_ubuntu-server-01_elastic_agent_rollout.md`. Runs **OWASP Juice Shop live** on port 3000 (HTTP-confirmed, not a leftover — it's up right now).

3.6 Security Onion — SOC platform

Security Onion 3.1.0 standalone. IP `.30`, SSH `security-onion / socadmin` (key auth works). OS timezone intentionally UTC (web UI already renders local time). Own internal firewall — hostgroups gate portgroups, membership in one grants nothing in another (root cause of the DC01 outage, §7.6; reference: §5.4). Bundles Suricata, Zeek, Elasticsearch, Kibana, Fleet, Hunt.

3.7 Kali Linux — Red Team workstation

VM `ATTACK-Kali` (leading space). IP `.50` , SSH `kali / blue1` (key auth works). Vulnerability scanning, exploitation, network analysis — authorized use inside this isolated lab only.

3.8 Metasploitable2 — vulnerable training target

VM `Target-Metasploitable2` . IP `.70` . Stock, unmodified port fingerprint confirmed (FTP/SSH/Telnet/SMTP/DNS/HTTP/rpcbind/SMB/rexec-rlogin-rsh/Java-RMI/ingreslock/NFS/ProFTPD-backdoor/MySQL/distccd/PostgreSQL/VNC/X11/UnrealIRCd/AJP/Tomcat). No ACPI shutdown support — `lab-stop.sh` force-stops it after 60s (expected).

4. Architecture, security & AI rules

Why these technologies were chosen

Decision	Choice	Core reason
Host OS	Bazzite Linux	Modern Linux + NVIDIA + gaming + virtualization on one machine
Virtualization	KVM/QEMU/libvirt/virt-manager	Native Linux, full control, deeper understanding than a turnkey hypervisor
Firewall	OPNsense	Realistic enterprise firewall/routing/DHCP/DNS practice
SOC platform	Security Onion	Complete SIEM+IDS+NSM bundle, superseding a hand-rolled stack
Identity	Windows Server AD (DC01)	Realistic detection/attack scenarios (logins, privilege escalation, lateral movement)
Version control	Git	History, recovery points, professional change tracking
Secrets storage	Separate <code>Secure/SOC-Secure.img</code> , never in Git	Docs can be shared; secrets must never enter Git history
AI assistance	ChatGPT (planning) + Claude Code (implementation)	Support tools, not decision-makers — bound by the rules below

AI rules (binding on any AI assistant working on this project)

Merged from `AI_ACCESS_POLICY.md` , `PROJECT_RULES.md` , `CLAUDE.md` . Claude Code acts as a **junior SOC engineer and documentation assistant**, not an autonomous administrator — documentation is treated as more reliable than its own memory; source of truth: `README.md` ,

`PROJECT_RULES.md` , `AI_ACCESS_POLICY.md` , `docs/INDEX.md` , `docs/guides/` ,
`docs/troubleshooting/` , `docs/decisions/` .

Allowed: read/analyze documentation and configs, explain commands, suggest improvements, write/maintain documentation, troubleshoot (research, diagnosis, proposing fixes). AI assists — it doesn't replace administrator decisions.

Never: store or write down passwords/secrets/API keys/private keys/tokens anywhere; commit credentials; modify firewall rules or critical infrastructure without approval; delete or destructively change anything without confirmation; ignore previous documented decisions.

Credentials: never in Git or markdown, not even temporarily. Live only in `Secure/SOC-Secure.img` , opened manually by the administrator — an assistant never touches it. The user enters passwords themselves when needed.

Infrastructure change procedure: explain what will change → explain risk/impact → snapshot/backup → get explicit confirmation → execute the smallest sufficient change → document the result.

Troubleshooting method: check existing docs → identify the affected system → review prior troubleshooting cases → explain root cause → apply the smallest safe fix → update documentation.

General behavior: respect prior decisions, preserve existing architecture, ask before destructive (even reversible) actions, document discoveries as they happen, prefer safe/reversible solutions, explain commands and verify results rather than assuming success.

Read-only investigation is in-scope without prior confirmation — e.g. using `virsh` , ARP tables, unauthenticated port probes, or already-working passwordless SSH to verify facts for documentation. Anything that *changes* the lab (firewall, AD objects, service state) still needs the full change procedure above.

Final principle: documented, reproducible, secure, understandable — documentation first, automation second, security always.

Core security principles

Least privilege · defense in depth (OPNsense + Security Onion hostgroups + host controls) · network isolation (192.168.50.0/24 is isolated because it hosts attack tooling and deliberately vulnerable machines) · documentation before/with automation · backups before changes · monitoring as the foundation, not an add-on · all offensive testing stays inside the lab boundary.

Incident response phases

Identification → Analysis → Containment → Recovery → Lessons learned. Operationalized in the runbook — [§6.2](#).

5. Operational guides

5.1 Starting/stopping the lab

Preferred: **Pentest Lab Start / Stop** desktop launchers (all 7 VMs, correct order). Command line:

```
~/Homelab/scripts/lab-start.sh
~/Homelab/scripts/lab-stop.sh
```

Single VM: `virsh -c qemu:///system start|shutdown <vm-name>` — note the leading space in `ATTACK-Kali` ([§9](#)). `lab-start.sh` is idempotent; `lab-stop.sh` tries graceful ACPI shutdown first (60s/180s timeout) before force-killing.

5.2 Connecting to systems

SSH All Machines launcher opens one Konsole tab per machine, or directly:

```
ssh opnsense
ssh dc01
ssh security-onion
ssh kali
ssh ubuntu-server
ssh win11-01
```

Homelab VM Manager opens virt-manager on `qemu:///system`. **Key status:** `kali`, `security-onion`, `dc01`, `ubuntu-server`, and `win11-01` are all passwordless; only `opnsense` still prompts (password-only by design, not a regression — see [§9](#)).

5.3 Security Onion / Kibana / Fleet in the browser

Security Onion Operator launcher opens a persistent Playwright/Chromium profile with every key page as a tab (Overview, Hunt, Detections, Cases, Grid, Kibana, Fleet), driven via a local CDP port (never exposed beyond localhost). First run waits for manual login to both Security Onion and Kibana/Fleet (two separate logins); session persists in `browser/profile/` (never in Git). Everything built so far is **read-only** — nothing can modify Fleet, detection rules, or Security Onion config.

Fleet status report: `~/Homelab/scripts/soc-web-audit.sh` → Markdown report in `browser/artifacts/`.

 **Recorded incident:** a script once printed a live session cookie into terminal output via a raw `context.request` call. No external leak, but fixed immediately (daemon restarted, session invalidated) and all API calls now route through `fetchJsonInPage()`, which runs inside the authenticated page so cookies never pass through this project's own code. Don't reintroduce raw out-of-page HTTP calls without the same care.

5.4 Network ports & firewall hostgroups reference

Security Onion's **own** firewall (separate from OPNsense) — reconstructed from `/opt/so/saltstack/default/salt/firewall/{defaults,soc_firewall}.yaml`.

Portgroup	Port(s)	Purpose
elastic_agent_control	8220	Fleet Server checkin
elastic_agent_data	5055	Data ingest
elastic_agent_update	8443	Agent updates
endgame	3765	Elastic Defend/Endpoint output
beats_5044 / 5644 (SSL) / 5066 / 5056	—	Logstash beats input
kibana	5601	Kibana web UI
elasticsearch_rest	9200	Elasticsearch REST

Key hostgroups: `analyst` (web UI/Kibana users) · `elastic_agent_endpoint` (Fleet checkin) · `beats_endpoint` / `_ssl` (log shipping) · `endgame` (Elastic Defend) · `manager` / `receiver` / `sensor` / `standalone` (grid roles) · `desktop` (admin workstations, like `analyst` plus extras).

A normal endpoint with Elastic Agent needs *all three* of `elastic_agent_endpoint`, `beats_endpoint`, and `endgame` — `analyst` alone is not enough; one hostgroup grants nothing in another.

```
sudo so-firewall includehost <hostgroup> <ip>
sudo so-firewall apply          # additive; so-firewall removehost <ip> to strip
entirely
```

Read-only inspection: `cat /opt/so/log/so-firewall.log`, `ss -tlnp`. From another host: `Test-NetConnection -ComputerName 192.168.50.30 -Port 8220`.

5.5 Desktop launchers

Launcher	Purpose
Pentest Lab Start / Stop	Start/stop all 7 VMs in order, via Konsole (<code>--hold</code>)
SSH Alle Machines	One Konsole window, independent tab per machine
Homelab VM Manager	virt-manager on <code>qemu:///system</code>
Security Onion Operator	Persistent browser session — §5.3

Bazzite has no `gnome-terminal` — all terminal launchers use Konsole specifically.

SSH config (`~/.ssh/config`):

Alias	Host	User	Status
opnsense	.1	root	Password-only, currently not even password-prompting reliably
dc01	.10	Administrator	Key auth works
security-onion	.30	socadmin	Key auth works
kali	.50	blue1	Key auth works
ubuntu-server	.40	sysadmin	Key auth confirmed working (2026-07-14)
win11-01	.20	pentest\administrator	Added 2026-07-14, key auth confirmed working (see docs/troubleshooting/09_win11-01_ssh_access.md)

5.6 Quick reference

```
# Sysmon test event on DC01, then check Hunt (Last 15 Minutes):
# host.name:"dc01" AND event.dataset:"windows.sysmon_operational"
Start-Process -FilePath 'cmd.exe' -ArgumentList '/c echo test' -Wait

~/Homelab/scripts/soc-mirror.sh --status # mirroring status
ssh security-onion "tail -20 /opt/so/log/so-firewall.log"
```

6. Detection use cases & incident response

6.1 What this SOC should detect

 present (works once traffic/logs arrive) ·
  should work via default rulesets, not yet deliberately confirmed ·
  needs custom rules, not built.

Detection	Status
ICMP ping sweeps / TCP scans (SYN/FIN/NULL/XMAS) / UDP scans	✅ Confirmed 2026-07-15 — see Metasploitable2 Tier 1 scan below
OS fingerprinting / banner grabbing	✅ Confirmed 2026-07-15 — see Metasploitable2 Tier 1 scan below
SSH / FTP / SMB brute force	⚠️ (SMB relevant especially toward DC01)
AD / LDAP / SMB enumeration	✅ Confirmed 2026-07-15 — see "DC01 Tier 1 AD enumeration" below. Deliberately a separate row from brute force: enumeration (single null-session/anonymous-bind probes reading directory/share structure) and brute force (repeated authentication attempts against credentials) are different behaviors with different signatures, and shouldn't be collapsed into one row.
Suspicious DNS requests	⚠️ (Zeek <code>dns.log</code> ; Sysmon event 22 on DC01)
Known exploit signatures / reverse shells / Metasploit indicators	✅ Confirmed 2026-07-15, with an important nuance — see "Metasploitable2 Tier 2 exploitation" below
SQLi / command injection / directory traversal	⚠️ (Suricata web-attack rulesets)

Metasploitable2 Tier 1 scan — TCP scans + OS fingerprinting/banner grabbing, confirmed 2026-07-15:

- **Command:** `nmap -sV -sC -p- 192.168.50.70` (full 65535 TCP ports, service/version detection, default NSE scripts).
- **Source:** ATTACK-Kali (`192.168.50.50` , run via SSH from the Bazzite host). **Target:** Metasploitable2 (`192.168.50.70`).
- **Time window:** `2026-07-15 01:50:48Z – 01:53:20Z` (scan itself: `01:50:54Z – 01:53:10Z` , 136.31s per nmap's own summary).
- **Result:** 30 open ports found (vsftpd 2.3.4 with anonymous login, telnet, Postfix smtpd, BIND 9.4.2, Apache 2.2.8, unauthenticated NFS/RPC, Samba 3.0.20, rexecd, a literal "Metasploitable root shell" bindshell on 1524, MySQL 5.0.51a, distccd, PostgreSQL 8.3, VNC, UnrealIRCd, Tomcat, Ruby DRb) — matches this target's known stock fingerprint.
- **Hunt evidence, same source/dest/window:**
 - `event.dataset:"zeek.conn" / zeek.weird : 65,801 events` — confirms full network-level visibility of the scan traffic (query: `source.ip:"192.168.50.50" AND destination.ip:"192.168.50.70" AND @timestamp:[2026-07-15T01:50:40.000Z TO 2026-07-15T01:53:20.000Z]`).
 - `event.module:"suricata"` , same source/dest/window: **172 real alerts**, not raw traffic — signature-based detections firing on the actual scan. Distinct signatures observed: `ET SCAN Potential SSH Scan OUTBOUND` (medium), `ET SCAN Suspicious inbound to PostgreSQL port 5432` (medium), `ET SCAN Suspicious inbound to mySQL port 3306` (medium), `GPL DNS named version attempt` (medium — fired by `-sC`'s `dns-nsid` script, a direct banner-grab signature), `GPL NETBIOS SMB-DS IPC$ share access` (low — fired by the `SMB/ smb-os-discovery` enumeration scripts).

- **Why this flips both rows:** ET SCAN Potential SSH Scan OUTBOUND and the PostgreSQL/MySQL "suspicious inbound" signatures are pattern-based scan detections that fired directly on this `-p-` full-port sweep — confirms **TCP scan detection** works, not just default-ruleset assumption. GPL DNS named version attempt and the SMB share-access/ smb-os-discovery signature are specifically banner-grab/service-fingerprinting detections, triggered by nmap's own `-sC` scripts querying BIND's version and enumerating SMB — confirms **OS fingerprinting/banner grabbing detection** works, not just passive Zeek visibility as previously assumed.

Host-based (Sysmon on DC01): Process Create (1), Network Connection (3), File Create (11), DNS Query (22) — all  confirmed with test events. Other event types (registry, image load) —  covered by config, not individually tested. Elastic Defend also produces its own detections since the 2026-07-13 firewall fix.

How to test (Purple Team style): run a harmless version of the technique from Kali, check Security Onion Hunt/Detections. Example:

```
for i in 1 2 3 4 5; do ssh -o ConnectTimeout=3 -o BatchMode=yes
nope@192.168.50.40; done
# Hunt: destination.ip:"192.168.50.40" AND event.dataset:*ssh*
```

6.2 Incident response runbook

Fixed procedure for any alert — real, test, or exercise:

1. **Confirm** — Security Onion Alerts/Detections: which rule, source/destination, when (SO shows local time, +02:00 summer), how often.
2. **Context** — Hunt: `source.ip:"<IP>" OR destination.ip:"<IP>"` . Known lab system ([§9](#))? Deliberate test running (`docs/daily/`)? Anything else at that time?
3. **Host check** (if SSH-reachable) — logs/processes, Sysmon events around the timestamp. Investigating is always fine independently; changing/restarting/isolating needs a quick explanation first unless already agreed.
4. **Record** — write up in `docs/daily/YYYY-MM-DD/rapport.md` regardless of outcome; genuine new technical problems also get a `docs/troubleshooting/` doc.
5. **If real** — don't panic (lab is isolated); consider isolating via `virsh` (explain first); preserve evidence before cleanup; document as step 4.

Quick Hunt queries: `host.name:"<lowercase-name>" . ... AND event.module:"windows" . ... AND event.dataset:"windows.sysmon_operational" . source.ip:"<ip>" OR destination.ip:"<ip>"` .  Hostnames index lowercase (`dc01`), even though Fleet displays capitalized.

6.3 Detection validation plan

In progress — [§12](#) test session started 2026-07-15 (Phase 3 of the project roadmap). Every  in [§6.1](#) gets tested here rather than in a separate pass: Tier 1 (recon) covers scans/sweeps/fingerprinting/DNS; Tier 2 (Metasploitable2/Juice Shop exploitation) covers brute force,

exploit signatures, reverse shells, Meterpreter, SQLi/injection; Sysmon "other" event IDs get tested on DC01 or WIN11-01 once it's a target. Method: run technique → check Hunt → flip 🚩 to ✅/❌ with the query and evidence used. Also planned: one full runbook dry-run (§6.2) against a real alert from that pass, logged in docs/daily/ .

Tier 1 progress:

Scenario	Status	Evidence
Full port/service scan — Metasploitable2	✅ Done, 2026-07-15	See §6.1 "Metasploitable2 Tier 1 scan" — 172 Suricata alerts, 65,801 Zeek events, flipped 2 rows to ✅
Web app recon — Juice Shop (nikto , gobuster / ffuf)	✅ Done, 2026-07-15	See "Juice Shop Tier 1 web recon" below — 3,637 Suricata alerts, no §6.1 row flipped outright (see notes)
AD/domain enumeration — DC01 (enum4linux-ng , netexec smb , anonymous LDAP)	✅ Done, 2026-07-15, read-only only	See "DC01 Tier 1 AD enumeration" below — 3 Suricata alerts incl. a high-severity anonymous-LDAP-bind signature; flipped a new §6.1 row, "AD / LDAP / SMB enumeration," to ✅
Network sweep (.0/24)	✅ Already done (nmap -sn , confirmed known IPs only, per §12 "Preparation already done")	—

Juice Shop Tier 1 web recon — confirmed 2026-07-15:

- **Commands:** nikto -h http://192.168.50.40:3000 (8,907 requests, 95s, 25 items reported), then gobuster dir -u http://192.168.50.40:3000 -w /usr/share/wordlists/dirb/common.txt --exclude-length 9903 (the Juice Shop SPA returns HTTP 200 for any unmatched path — length-excluded to filter that wildcard).
- **Source:** ATTACK-Kali (192.168.50.50). **Target:** ubuntu-server-01 / Juice Shop (192.168.50.40:3000).
- **Time window:** 2026-07-15 01:58:31Z — 02:01:09Z (nikto 01:58:34 — 02:00:09Z ; gobuster 02:00:14 — 02:01:09Z , back-to-back, both timestamped independently).
- **Nikto findings:** exposed .bash_history / .sh_history / .htpasswd -shaped paths, several *.json endpoints ("might be interesting"), a x-recruiting header, missing CSP/referrer-policy/HSTS headers — informational web-hygiene findings, not a Suricata-detection test in themselves.
- **Gobuster findings:** real endpoints confirmed — /api , /rest* , /assets , /media , /promotion , /ftp , /robots.txt , plus an oddity worth a separate look later (/Video and /video both returning a 2.2 MB body).
- **Hunt evidence, same source/dest/window, port 3000: 4,557 total events, of which 3,637 are real Suricata alerts** (event.module:"suricata"), not raw traffic. Distinct signatures: ET EXPLOIT QNAP Shellshock CVE-2014-6271 (**high**, Attempted Administrator Privilege Gain — nikto's Shellshock-probe header), GPL WEB_SERVER iisadmin access (**high**, Web Application Attack), ET WEB_SERVER WEB-PHP phpinfo access , GPL WEB_SERVER global.asa access ,

GPL WEB_SERVER printenv access , GPL WEB_SERVER /~root access , ET WEB_SERVER Possible Cherokee Web Server GET AUX Request DoS Attempt , ET INFO Proxy TRACE Request - inbound , ET INFO Dotted Quad Host DLL Request .

- **Why no §6.1 row is flipped outright:** these signatures are mostly nikto's own vulnerability-probe requests and gobuster's wordlist entries coincidentally matching legacy CGI/PHP/IIS signature paths — real detections, but not a clean match for any existing §6.1 row. The Shellshock signature is the closest (ET EXPLOIT . . .), but "Known exploit signatures / reverse shells / Metasploit indicators" also explicitly covers reverse-shell/Meterpreter indicators, which are untested here (Tier 2 scope, not yet approved) — flipping that row fully would overclaim based on this evidence alone. Recorded as strong supporting evidence for that row once Tier 2 actually tests exploitation.

DC01 Tier 1 AD enumeration — confirmed 2026-07-15, read-only only:

- **Commands:** `enum4linux-ng -A 192.168.50.10` (full read-only sweep: SMB dialects, null-session domain info, RPC users/groups/shares/policy/printers), `netexec smb 192.168.50.10 --shares` (anonymous share enum), `ldapsearch -x -H ldap://192.168.50.10 -b "dc=pentest,dc=lab" -s base` (anonymous LDAP bind). No credentials used anywhere, no write/modify operation attempted.
- **Source:** ATTACK-Kali (192.168.50.50). **Target:** DC01 (192.168.50.10).
- **Time window:** 2026-07-15 02:04:46Z – 02:05:05Z .
- **AD-side result — DC01 correctly rejected enumeration:** null-session SMB connected and revealed only basic domain identity (NetBIOS name DC01 , domain PENTEST / pentest.lab , domain SID) — every actual enumeration call (`querydispinfo` , `enumdomusers` , `enumalsgroups` , `enumdomgroups` , share list, policy, printers) returned STATUS_ACCESS_DENIED . `netexec` confirmed null-auth accepted at the protocol level but got STATUS_USER_SESSION_DELETED on share enumeration. Anonymous LDAP bind was rejected outright (`Operations error` — "a successful bind must be completed"). This is a real, confirmed security posture, not assumed.
- **Hunt evidence, same source/dest/window:** 76 total events; **3 real Suricata alerts** (`event.module:"suricata"`): ET INFO NTLM Session Setup Request - Auth (low), ET INFO NTLM Session Setup Request - Negotiate (low), and **ET INFO Anonymous LDAPv3 Bind Request Outbound (high, "Potential Corporate Privacy Violation")** — a purpose-built signature that fired precisely on the anonymous LDAP bind attempt.
- **New §6.1 row added, 2026-07-15: "AD / LDAP / SMB enumeration," status .** None of the original 6 rows was a clean match — closest was "SSH/FTP/SMB brute force," but that implies repeated authentication attempts against credentials, not single null-session/anonymous-bind probes reading directory/share structure. These are different behaviors with different signatures (ET INFO Anonymous LDAPv3 Bind Request Outbound / NTLM Session Setup vs. e.g. repeated-failed-logon signatures) and belong in separate rows, not collapsed into one — Joost's explicit decision, not a unilateral restructure.

Tier 1 complete as of 2026-07-15 — all three planned scenarios (Metasploitable2 full-port scan, Juice Shop web recon, DC01 read-only AD enumeration) executed and validated in Hunt with direct evidence.

Tier 2 explicitly approved by Joost, 2026-07-15 ("test tier 2 met mijn toestemming") — execution started the same day, one scenario at a time per [[feedback_phase3_attack_testing_discipline]].

Scenario	Status	Evidence
vsftpd 2.3.4 backdoor — Metasploitable2	✅ Done, 2026-07-15	See "Metasploitable2 Tier 2 exploitation" below — real root RCE achieved, nuanced detection result
Samba/NFS/RMI — Metasploitable2	⌚ Not yet run	—
OWASP Top 10 — Juice Shop	⌚ Not yet run	—
UnrealIRCd backdoor — Metasploitable2	⌚ Not yet run	—

Metasploitable2 Tier 2 exploitation — vsftpd 2.3.4 backdoor (CVE-2011-2523), confirmed 2026-07-15:

- **Technique:** the classic vsftpd 2.3.4 malicious-backdoor RCE — repeats the exact exploit already noted in §8 as done once before, outside this documentation trail. First attempted via `msfconsole` (`exploit/unix/ftp/vsftpd_234_backdoor`), which triggered the backdoor correctly (its own console output confirmed `Backdoor has been spawned!`) but didn't cleanly hand control to a Metasploit session due to a payload/session-handler mismatch in this Metasploit version. Reading the module's own source (`/usr/share/metasploit-framework/modules/exploits/unix/ftp/vsftpd_234_backdoor.rb`) showed the exact underlying mechanism: send `USER <random>:` (a smiley in the FTP username is the trigger), then `PASS <random>` , then connect directly to TCP/6200 for a root shell. Reimplemented that sequence directly in a small Python script for a clean, transparent demonstration, run three times total (two earlier via the `msfconsole` attempts, one final clean manual run).
- **Source:** ATTACK-Kali (`192.168.50.50`). **Target:** Metasploitable2 (`192.168.50.70`).
- **Result — real root code execution, all three attempts:** `id` → `uid=0(root) gid=0(root)` , `whoami` → `root` , `uname -a` → `Linux metasploitable 2.6.24-16-server ... i686 GNU/Linux` , `hostname` → `metasploitable` . No files modified, no persistence installed, no pivoting attempted — read-only recon commands only, consistent with this project's lab-exploitation scope.
- **Detection result — nuanced, both halves independently confirmed in Hunt:**
 - **The exploit trigger itself (FTP `USER x:`) / `PASS x` , then the connection to 6200) was NOT detected.** `event.module:"suricata"` for `source.ip:"192.168.50.50"` AND `destination.ip:"192.168.50.70"` in the exact exploit windows returns zero hits — only Zeek `conn / weird ( active_connection_reuse` , an artefact of the rapid back-to-back connections) logged the traffic at all. No ET/GPL signature exists in this ruleset for the vsftpd-backdoor trigger sequence specifically.
 - **The post-exploitation confirmation WAS detected, all three times.** `GPL` `ATTACK_RESPONSE id check` returned `root` fired on the *reverse-direction* traffic (`source.ip:"192.168.50.70"` port 6200 → `destination.ip:"192.168.50.50"` , medium severity, category "Potentially Bad Traffic") at `11:48:21.761Z` , `11:48:50.650Z` , and `11:53:46.630Z` — each timestamp lines up exactly with an `id` command's root-output response crossing the wire. This is a generic content-inspection signature (it looks for `uid=0(root)` appearing in any TCP response, not a vsftpd-specific rule) — it would fire on *any* technique that causes root `id` output to traverse the network, not just this one.

- **Why this flips the §6.1 row, with a caveat:** the row asks about "known exploit signatures / reverse shells / Metasploit indicators" broadly, and a real, reproducible signature did fire on real successful exploitation, three times, with second-level timestamp precision. But it's honest to note this is **detection of the symptom (a root shell responding), not the specific exploit** — a more targeted attacker who avoided ever running `id` (or similar output-leaking commands) over that particular connection would not have tripped this signature. This is a genuine, useful, but narrow detection — not a dedicated vsftpd-backdoor rule. Recorded as a real capability, not overclaimed as complete coverage.
- **SOC Alarmdashboard cross-check:** the same `GPL ATTACK_RESPONSE id check returned root` alerts also appeared in the live alarm dashboard, bucketed as `SCAN` (via the generic "potentially bad traffic" category-text fallback) — arguably a better fit for `EXPLOIT` given it's literally a post-compromise confirmation, not reconnaissance. Not changed in this pass to avoid re-iterating the categorizer without being asked; flagged here for a future categorization refinement.

Tier 2 continues with the remaining three scenarios (Samba/NFS/RMI, Juice Shop OWASP Top 10, UnrealIRCd backdoor), one at a time. **Tier 3 remains explicitly out of scope without new, separate explicit approval.**

7. Troubleshooting history

Real problems from building this lab. Full evidence/commands/rollback steps live in `docs/troubleshooting/0N_*.md` — this is the summary.

#	Problem	Root cause	Fix / lesson
7.1	QEMU guest agent not configured	Not installed/enabled in-guest, channel not configured	Install + enable + configure channel. Lesson: a running VM ≠ working guest integration — test explicitly.
7.2	virt-manager VM setup	Storage/display quirks	Standardized on qcow2; SPICE issues traced to guest tooling, not hardware. Snapshot before risky changes.
7.3	OPNsense network debugging	Firewall issues masquerading as app/DNS problems	Verify bottom-up: interface → IP → gateway → DNS → firewall. DNS misconfig was repeatedly the real AD root cause.
7.4	Security Onion installation	Under-provisioned resources	SOC platforms need materially more RAM/CPU/storage than an ordinary server. A SOC is only as useful as the telemetry reaching it.
7.5	DC01 AD setup	—	DNS is the foundation AD depends on. Snapshot before promoting to DC (hard to cleanly undo).
7.7	<code>soc-mirror.sh</code> libvirtd deadlock	Synchronous <code>virsh</code> callback from inside a libvirt qemu hook could hang libvirtd	Detached the reconciler via <code>systemd-run --no-block --collect</code> .
7.8	libvirt/OPNsense DHCP conflict (2026-07-09)	libvirt's own DHCP collided with OPNsense's; <code>virbr10</code> claimed OPNsense's own IP	Stripped libvirt's <code><ip></code> / <code><dhcp></code> from the network definition — OPNsense is sole DHCP/DNS/gateway authority. Host got a separate mgmt IP (.254) restored at boot.

7.6 DC01 Fleet health & Sysmon telemetry (2026-07-13)

The largest case — three independent, stacked root causes behind one symptom ("DC01 Offline in Fleet"). Recurring: a 2026-07-10 "fix" was a non-persistent `iptables` rule that any `so-firewall apply /reboot` silently discarded.

- Firewall hostgroups** — DC01 was only ever added to `analyst`, never `elastic_agent_endpoint` / `beats_endpoint` / `endgame` (each hostgroup only grants its own ports — see §5.4). Fixed with `so-firewall includehost` for all three (additive), `so-firewall apply`.
- ~9h clock skew** — `vmictimesync` was fighting NTP at every boot (DC01, as PDC Emulator, doesn't behave like an ordinary NTP client). Disabled the service, forced `w32tm /resync`, added a boot-time scheduled task. Side effect: the skew had poisoned Logstash's dead-letter queue with rejected future-timestamped events — explains why `HEALTHY`-reporting agents still showed no data. **Lesson: agent-HEALTHY proves shipping, not acceptance.**
- Sysmon never installed** — Fleet policy expected `windows.sysmon_operational` data that could never arrive. Installed Sysmon 15.21 + SwiftOnSecurity config.

Verified: survives Elastic Agent restart, two DC01 reboots, and a full Security Onion reboot. Rollback procedures (firewall/clock/Sysmon/timezone) kept in

`docs/troubleshooting/06_dc01_fleet_health_and_sysmon.md` , not repeated here.

General lessons: a "Healthy" status can hide multiple stacked root causes — validate across a real reboot, not just current state. Per-hostgroup firewall reachability doesn't generalize across ports. PDC Emulators aren't ordinary NTP clients.

8. Project timeline

Period	What happened
Before 07-05	Fortress Bazzite: virt-manager, Kali/Debian, Ubuntu Server, Metasploitable2, Docker, Juice Shop, Burp Suite, Metasploit set up. Design doc planning a Suricata+Zeek+Wazuh stack.
07-03 – 07-07	Kali/Metasploitable2 VMs up, <code>pentest-lab</code> network built. Metasploitable2 enumerated and exploited (<code>vsftpd_234_backdoor</code> → Meterpreter). Juice Shop deployed, first vuln found via Burp Suite. AX210 WiFi passed through to Kali. Host-level Suricata installed as a side experiment (<code>/var/home/Joost/FORTRESS/</code> , separate from this repo). OPNsense installed 07-07 (redone with ZFS after a failed first attempt); WAN/DHCP issues left open.
07-08	Transition to "SOC Homelab" naming (⚠ reconstructed from limited evidence).
07-09	First git commit. Fixed libvirt/OPNsense DHCP conflict (§7.8); Kea DHCP reservations created; launchers fixed for KDE/Konsole.
07-10	DC01 built (AD DS + DNS, promoted, users created), enrolled in Fleet. First Fleet issue "fixed" non-persistently (see §7.6).
07-11	Full documentation structure created (README, rules, AI policy, network/server/AD docs, troubleshooting history).
07-12	Secret scrubbed from git history. Desktop launchers rebuilt properly. Traffic mirroring rewritten (deadlock fix). Browser automation for Security Onion built.
07-13	DC01 Fleet outage fully resolved. Live lab verification pass (IPs, AD structure, hardware). Documentation substantially expanded, including this master document, its AI-rules and attack-scope sections. Commit made, not pushed (standing instruction: no push without explicit permission).

9. Asset inventory

Full detail: `docs/ASSET_INVENTORY.md` .

Physical host

Property	Value
OS	Bazzite Linux
CPU	Intel Core i9-11900K (8c/1s)
GPU	NVIDIA GeForce RTX 3090
RAM	62 GiB
WiFi	Intel AX210 (Wi-Fi 6E)
Virtualization	KVM, QEMU, libvirt, virt-manager

Virtual machines

VM name (virsh)	IP	SSH	OS / role
OPNsense-FW	.1	opnsense /root	Firewall/gateway
DC01	.10	dc01 /Administrator	Windows Server 2022, AD DC, domain pentest.lab
WIN11-01	.20	win11-01 / pentest\administrator	Windows 11, domain-joined as DESKTOP-EFKB8GQ
SOC-SecurityOnion	.30	security-onion /socadmin	Security Onion 3.1.0 standalone
ubuntu-server-01	.40	ubuntu-server /sysadmin (key auth works)	Linux server, live Juice Shop on :3000
ATTACK-Kali	.50	kali /blue1	Red Team workstation
Target-Metasploitable2	.70	none	Vulnerable target, stock fingerprint

Active Directory identity inventory

OU	Accounts
Domain Controllers	DC01
Admins	IT Admin 01 (not in Domain Admins)
AD-Users	soctest, Helpdesk 01, Employee 01, Manager 01, HR 01, Finance 01
Service-Accounts	SQL Service
Workstations / Servers	(empty)

Custom groups: SOC-Analysts (soctest), Helpdesk (no members). Domain Admins = built-in Administrator only.

Software versions

Security Onion 3.1.0 · Elastic Agent 9.3.3+build202604082258 · Sysmon 15.21 (schema 4.91, SwiftOnSecurity config) · Windows Server 2022 Standard Evaluation (DC01).

Access methods (no passwords stored here)

SSH keys (`~/.ssh/config`) — passwordless for dc01/security-onion/kali/ubuntu-server/win11-01; password-only by design for opnsense. Browser session (Playwright, dedicated profile) for Security Onion/Kibana/Fleet, and a separate one for OPNsense (`browser/launch-opnsense-daemon.mjs` , port 9333). `virsh -c qemu:///system` for VM management.

Open items

- `opnsense` SSH password-only by design (confirmed via audit, not a regression) — key auth could still be added for convenience if desired.
- OPNsense's own config-revision history is empty (no built-in rollback safety net) and it hasn't checked for firmware updates since install — see `docs/OPNSENSE_AUDIT_2026-07-13.md` .
- The `KALI` firewall alias in OPNsense still points at Kali's old IP (`.157`) — harmless (unused in any active rule) but worth cleaning up.
- ~~`ubuntu-server-01` SSH key login doesn't work~~ — resolved 2026-07-14, key auth confirmed working (user `sysadmin` , not `ubuntu`).
- WiFi PCI-passthrough to Kali not reconfirmed.
- WIN11-01's intended training purpose — see [§12](#).

10. Glossary

Active Directory (AD) — Microsoft's user/computer/permission management system ("domain"). Runs on DC01; domain `pentest.lab` .

Agent — a program collecting/forwarding data (logs, metrics) to a central system (see Elastic Agent).

Beats — Elastic's lightweight data shippers; run inside Elastic Agent in modern deployments.

CEST/CET — Central European (Summer) Time; UTC+1 winter, UTC+2 summer.

Dead Letter Queue (DLQ) — where Logstash puts documents it couldn't store in Elasticsearch. A growing counter signals a downstream problem even without a sender-side error.

Domain Controller (DC) — AD server managing users/computers. Here: DC01.

Elastic Agent — collects Windows Event Logs/Sysmon/metrics on DC01 and ships to Security Onion; multiple components, each with its own health status.

Elastic Defend — Elastic Agent's endpoint-security (EDR) component; own Fleet connection (port 3765, hostgroup `endgame`).

Elasticsearch — Security Onion's log database; always UTC-internal.

Firewall hostgroup/portgroup — [§5.4](#). Hostgroup = named IP group, portgroup = named port group; Security Onion's firewall links them.

Fleet / Fleet Server — Kibana/Elastic's agent-management layer; server side runs on Security Onion, port 8220.

Hunt — Security Onion's own search UI (like Kibana Discover).

KVM/QEMU/libvirt — this lab's virtualization stack (kernel tech / hardware emulation / management layer).

NTP — clock-sync protocol; DC01 uses `pool.ntp.org`.

PDC Emulator — AD role: normally the domain's authoritative time source. DC01 holds it (the lab's only DC).

Security Onion — this lab's SOC platform (Suricata, Zeek, Elasticsearch, Kibana, Fleet). VM `SOC-SecurityOnion`, .30.

SIEM — collects/correlates logs across sources to find incidents. Security Onion is this lab's SIEM.

Sigma rule — general-purpose log detection rule format (Suricata's equivalent for network traffic).

so-firewall — CLI for Security Onion's own firewall (hostgroup membership). [§5.4](#).

Sysmon — Sysinternals tool logging detailed Windows events far beyond the standard Event Log. On DC01 since 2026-07-13, SwiftOnSecurity config.

UTC — global time standard, no offset. Elasticsearch always stores UTC internally.

vmctimesync — hypervisor clock-sync service (Hyper-V-compatible, also active under QEMU/KVM). Disabled on DC01 — was overriding NTP at boot.

Zeek (formerly Bro) — records network activity (DNS, HTTP, TLS, connections); complements Suricata's attack-pattern focus.

11. Current project status / what's next

Full detail: `docs/PROJECT_STATUS.md`.

Done

Base infrastructure (all 7 VMs on `pentest-lab`) · event-driven traffic mirroring · AD operational · Security Onion operational (web UI, Kibana, Fleet, Hunt) · DC01 Healthy in Fleet, survives restarts/reboots · passwordless SSH to security-onion/kali/dc01 · six desktop launchers · read-only web-audit script · this documentation structure · live network/asset/AD verification pass (2026-07-13) · read-only OPNsense configuration audit (2026-07-13) · Elastic Agent on the Bazzite host itself, Healthy,

confirmed reboot-survival, plus a central health-check script covering it and all 7 lab VMs (2026-07-14) · SSH access to WIN11-01 (2026-07-14) — OpenSSH Server enabled, integrated into `~/.ssh/config` and `lab-ssh-all.sh / soc-health-check.sh`, giving it the same admin path as the other lab systems · Elastic Agent + Sysmon on WIN11-01 (2026-07-14) — Healthy in Fleet, Elastic Defend/osquery/winlog/Sysmon/metrics all confirmed, telemetry verified in Hunt; endpoint-monitoring priority 1 of `docs/ROADMAP_ENDPOINT_MONITORING.md` done · Elastic Agent on ubuntu-server-01 (2026-07-14) — Healthy in Fleet, log/metrics-only, telemetry verified in Hunt; priority 2 done · ubuntu-server-01's long-standing `.100` -drift DHCP bug root-caused and fixed (2026-07-14) — see `docs/troubleshooting/12_ubuntu-server-01_dhcp_reservation_fix.md` · endpoint-monitoring phase fully closed (2026-07-15) — Kali will deliberately not get an Elastic Agent, Joost's final decision, not deferred · Phase 3 Tier 1 (recon) fully validated in Hunt (2026-07-15), see [§6.3](#) · SOC Alarmdashboard (2026-07-15) — live local banner+sound alerting on the Bazzite host per attack type, see `docs/guides/alarm_dashboard.md`.

⚠ Open

Security Onion OS-level timezone still UTC (cosmetic, blocked on root scope) · AD structural gaps (empty Helpdesk group, IT Admin 01 not elevated, undifferentiated role accounts, empty Workstations / Servers OUs) — **to be deliberately fixed as part of §12**, not left as-is · the full §12 test pass, AD escalation-path build, and WIN11-01 cleanup are scoped and agreed but **not yet executed** — pre-change snapshots already taken (DC01 : 2026-07-13-pre-ad-escalation-path , WIN11-01 : 2026-07-13-pre-target-cleanup) · the Bazzite host's own Elastic Agent's log delivery is now verified end-to-end into Elasticsearch (see `docs/troubleshooting/08_bazzite_host_elastic_agent.md`), but not yet re-confirmed across a reboot cycle.

✗ Planned

Detection engineering confirmation ([§6.3](#), Tier 1 done, Tier 2/3 pending approval) · combined host/security/gaming dashboard · four-tier alerting (INFO/WARNING/HIGH/CRITICAL) with Discord/Telegram forwarding — **note:** a *different*, attack-type-categorized local dashboard now exists (`docs/guides/alarm_dashboard.md`), this severity-tier/external-forwarding design remains separately unbuilt · VLAN segmentation, separate mgmt/attack networks, additional Windows clients, honeypots, periodic Purple Team exercises.

Keeping this current

Every important change: update the specific source doc → `CHANGELOG.md` → `docs/PROJECT_STATUS.md` → a daily report in `docs/daily/` → this master document.

12. Attack scope: agreed Red Team test plan

Scope agreed with Joost 2026-07-13; execution deliberately deferred to a dedicated session.

Grounded in the live-verified lab state above — nothing generic. Everything stays inside `192.168.50.0/24`, launched from Kali, against only the lab's own systems.

Agreed order:

1. **Tier 1 + 2 together, first** — recon/enumeration plus exploitation of the intentionally-vulnerable targets (Metasploitable2, Juice Shop). Low risk, closes most of §6.1's ⚠️ rows.
2. **AD attack chain (Tier 3) — build the escalation path first.** The environment currently has none (IT Admin 01 isn't a Domain Admin). Before Kerberoasting/privilege-escalation practice: set an SPN on SQL Service, give IT Admin 01 real elevated rights. This is an infrastructure change — full change procedure (§4), its own explained/confirmed step, not bundled into the attack.
3. **WIN11-01 becomes a target too, after cleanup** — move DESKTOP-EFKB8GQ into OU=Workstations, general tidy-up, *then* deliberately loosen its firewall for lateral-movement testing (WIN11-01 → DC01).

Preparation already done: pre-change snapshots (DC01 : 2026-07-13-pre-ad-escalation-path, WIN11-01 : 2026-07-13-pre-target-cleanup). Tier 1 (recon) fully executed and validated 2026-07-15 (§6.3). Tier 2 explicitly approved by Joost 2026-07-15 and started the same day — vsftpd 2.3.4 backdoor exploitation against Metasploitable2 done (real root RCE, nuanced detection result, see §6.3), three scenarios remain. **AD changes and firewall changes (Tier 3) remain not executed**, still pending separate explicit approval.

Tier 1 — Reconnaissance & enumeration

Scenario	Target	Technique
Full port/service scan	Metasploitable2 (.70)	<code>nmap -sV -sC -p-</code>
Web app recon	Juice Shop (.40:3000)	<code>nikto</code> , <code>gobuster</code> / <code>ffuf</code>
AD/domain enumeration	DC01 (.10)	<code>enum4linux-ng</code> , <code>netexec smb</code> , anonymous LDAP
Network sweep	.0/24	<code>nmap -sn</code>

Tier 2 — Exploitation of intentionally-vulnerable targets

Scenario	Target	Technique	Status
vsftpd backdoor	Metasploitable2:21	<code>vsftpd_234_backdoor</code> (repeats the historical exploit, §8)	✅ Done 2026-07-15, root RCE confirmed, see §6.3
Samba/NFS/RMI	Metasploitable2	<code>usermap_script</code> , anonymous NFS mount, Java-RMI deserialization	⌚ Not yet run
OWASP Top 10	Juice Shop	SQLi, broken auth, IDOR, XSS — built-in scored challenges	⌚ Not yet run
UnrealIRCd backdoor	Metasploitable2:6667	Known backdoor trigger	⌚ Not yet run

Tier 3 — Active Directory attack chain (after the escalation-path build)

Scenario	Target	Technique
Password spraying	DC01	Common passwords, low attempts/account
Kerberoasting	DC01	TGS request for <code>SQL Service</code> , crack offline (needs the SPN)
AS-REP roasting	DC01	Needs an account with Kerberos pre-auth disabled (none exists yet)
Lateral movement	WIN11-01 → DC01	From an assumed foothold — worth testing against the current locked-down firewall first
Golden/Silver ticket	DC01	Capstone — needs prior privileged access

Execution checklist: Tier 1+2 → cross-check Hunt, flip §6.1 statuses → build the AD escalation path (explained/confirmed) → clean up WIN11-01, loosen its firewall → run Tier 3 + lateral movement → update §6.1/§11 with results.

Document index

Root: `README.md` , `LAB_OVERVIEW.md` , `PROJECT_RULES.md` , `AI_ACCESS_POLICY.md` , `NETWORK.md` , `SERVERS.md` , `SECURITY.md` , `ACTIVE_DIRECTORY.md` , `CHANGELOG.md` . `docs/` : `INDEX.md` , `ASSET_INVENTORY.md` , `GLOSSARY.md` , `PROJECT_STATUS.md` , `OPNSENSE_AUDIT_2026-07-13.md` , `ROADMAP_ENDPOINT_MONITORING.md` , `PHASE1_CLOSURE_SUMMARY.md` , `ROADMAP_OPNSENSE_LOGGING.md` (Phase 2A — Firewall + DHCP syslog forwarding, validated 2026-07-15), `ROADMAP_PHASE2B_DNS_QUERY_LOGGING.md` (Unbound query logging, deliberately deferred 2026-07-15), `decisions/*.md` , `guides/*.md` (12 files, incl. `alarm_dashboard.md`), `troubleshooting/01 – 12` , `chat_history/*.md` (6 files), `daily/*/` , `screenshots/` .

`pentest.lab - by Joost Hebly.md` (repo root) is an earlier, unmerged concatenation of the same sources, built for the portfolio deliverable — this document supersedes it as the synthesized reference.